


AI FROM CLOUD TO YOUR PC

Local AI instead of the cloud — more privacy



AI as a tool

- Convenience
- Fast
- Relatively reliable answer
- Example:
 - Contracts
 - fast lookup of information



But what is with your data





Importance of Privacy

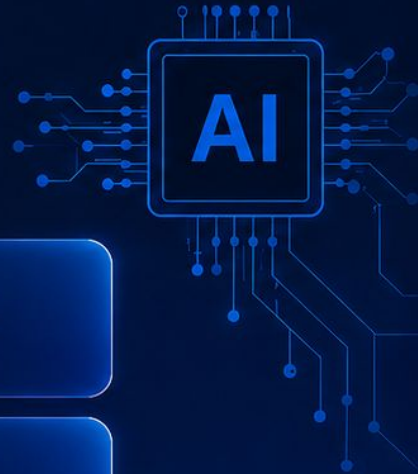
- Legal reasons for companies
- Your data is valuable
- You wouldn't share "secrets" with strangers

It's your data.

I just want you to think
about it for a moment.



The models



Name: Name of the Model (Qwen, DeepSeek, ...)



Format: Raw, GGUF, AWQ, etc.



Parameters: 0.6B, 7B, 35B-A7B, etc.

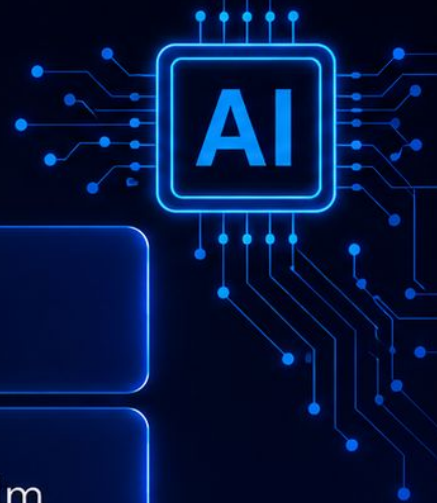


Quantization: Q2, Q4, Q8



Capabilities: think, tool-calling, image2image

The Backends



Easy, no brainer: Ollama



More Options, Command Line: llama.cpp, vllm



Cutting Edge: ktransformers

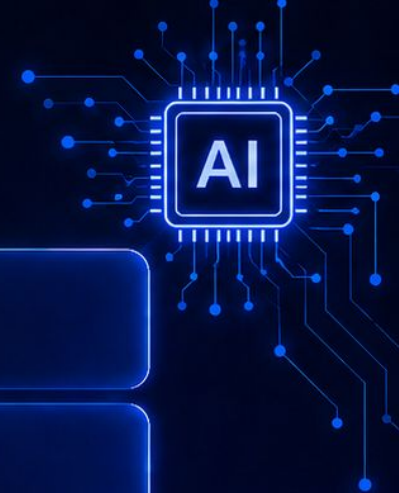


Raw: transformers



Specialized: ds4 for DeepSeek V4 Flash

Hardware



Size matters (VRAM/RAM) and speed



GPU is better — needs VRAM



CPU is also possible — needs RAM



Both combined is possible



Nvidia is better supported



Linux is preferred

>_ Example Ollama

```
# Package manager or download from the website
$ ollama --help

# Ollama serve & Launch ollama as server in the BG
$ ollama serve & # runs in background

# Ollama pull MODEL_NAME # download a model
$ ollama pull llama3.1:8b

# Ollama list # shows models
$ ollama list

# Ollama run MODEL_NAME # runs model in terminal
$ ollama run llama3.1:8b

# Easy
$ That's it. ■
```





Live Demo

What? A webinterface?



Free choice



Many clients



**Connect client with the URL
of the Ollama server output**



Limitations of Consumer Hardware



- Model: bigger, better, smarter

⚡ **but** ⚡



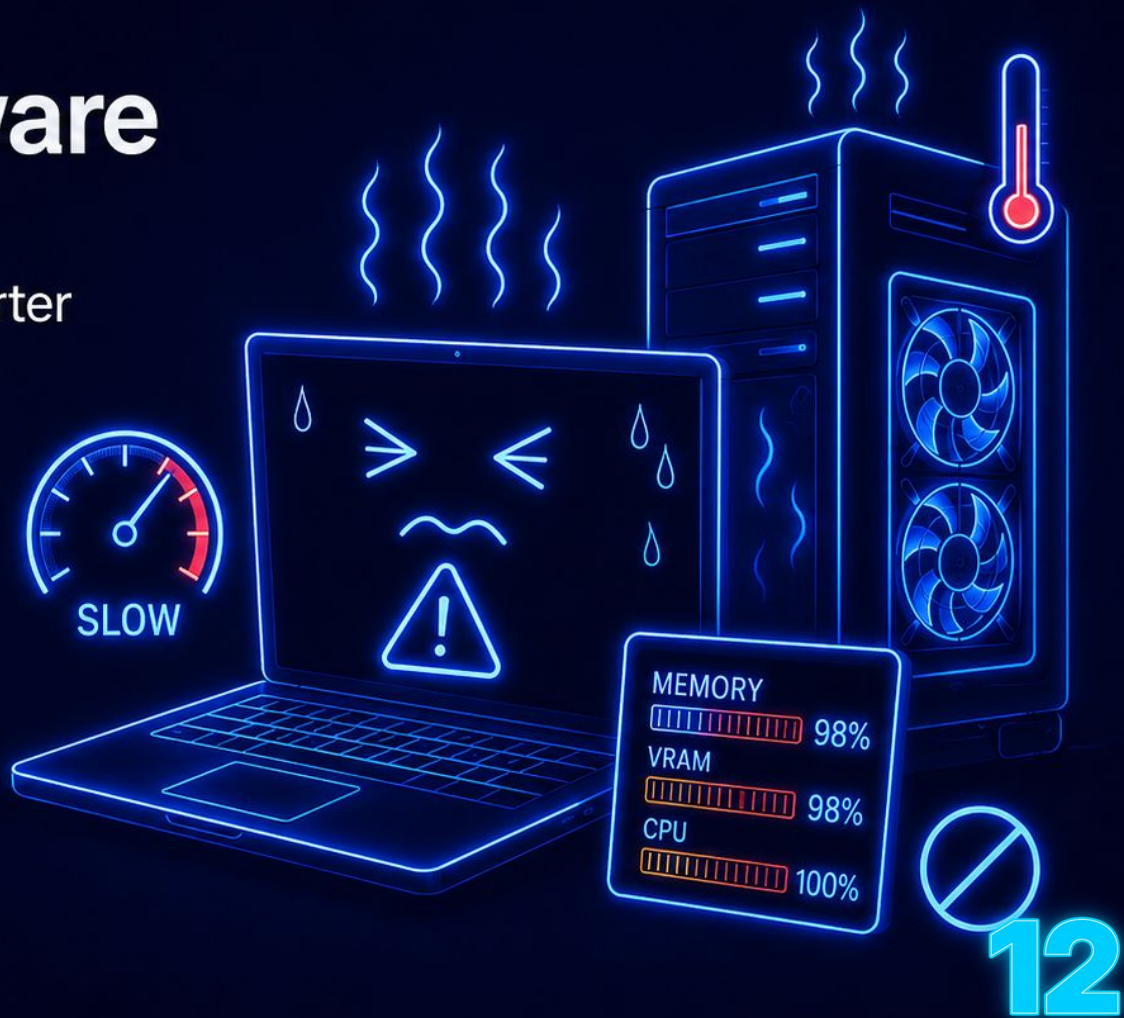
- Costs much



- Is slower



- Or not even runnable



>_

But it is possible
Happy limittesting

